Distributed Systems:
Crossing Intersection with Autonomous Vehicles

Antonio Toncetti
Gabriele Venturato

DMIF, University of Udine, Italy

September 11, 2019

Abstract

The aim of this project is to provide an implemented solution to the problem of autonomous vehicles crossing an intersection. Although the solution relies on some simplifications, it can be further elaborated to work in a real-case scenario.

The solution proposed in this report is meant to be more general and as modular as possible, in order for it to be possibly extended in a concrete situation.

Chapter 1

Introduction

The problem to solve is the one in which some autonomous vehicles have to cross an intersection without being involved in road accidents.

The idea is to solve the problem for a generic intersection. The autonomous vehicles can not rely on a central server, they have to cooperate with each other to cross the intersection by taking decisions which ensure a *fair* and *safe* policy. In particular there are two components in the proposed solution: vehicles and the environment. The latter is necessary in this context in order to simulate sensors that are usually inside autonomous vehicles which allows them to interact with the environment (e.g. proximity sensors, GPS, cameras, etc...). The system is *fault tolerant*, but neither byzantine processes nor cybersecurity hazards are taken in consideration, for simplification purpose.

The report starts by analysing the project: Chapter 2 is devoted to use cases, functional and non-functional requirements, and system assumptions too. Chapter 3 contains the description of the general architecture, and specific algorithms.

Following chapters aim to describe details of the implementation, and validation w.r.t. requirements.

Chapter 2

Analysis

This chapter describes in detail some fundamental assumptions about the system, as well as functional and non-functional requirements.

2.1 Use cases

2.1.1 Vehicle at an intersection

Brief Description

An Autonomous Vehicle (AV) is approaching, entering and leaving a non empty intersection. The use case begins with the AV approaching an intersection and ends with the AV having left the intersection in the desired direction.

Actors

- Autonomous Vehicles (AVs)
- Environment

Preconditions

- There is more than one AV at the intersection.
- AVs can be both active and inactive.
- The AV approaching the intersection is active.
- The AV knows the direction (exit) to take.

Scenarios

Main scenario

1. The AV is travelling along a road leading to an intersection.
2. The AV is approaching the intersection.
3. The AV stops at the entrance to the intersection.
4. The AV signals its intent to other AVs at the intersection.

5. The AV agrees with the other participants on how to solve the intersection.
6. The AV initiates the turn procedure towards the desired exit.
7. The AV is in the intersection travelling to an exit.
8. The AV leaves the intersection through an exit.
9. The AV signals that it has successfully left the intersection.
10. The AV continues to travel along its path.

Alternative scenario

1. 1, 2 The AV finds an active AV in front and gets in the queue.
2. 1, 2 The AV finds a faulty AV in front and overtakes him.
3. 6. 7 The AV finds a faulty AV and awaits for its removal.

Postconditions
At any point in its journey the AV can become inactive due to a mechanical failure or a software failure.

2.1.2   Vehicle mechanical failure

Brief Description
When an Autonomous Vehicle (AV) has a mechanical failure but the software is still working properly. The use case begins when the failure happens, and ends with the recovery of a normal situation.

Actors

- Autonomous Vehicle (AV)
- Environment

Preconditions

- There is at least one AV with a mechanical fault.
- There are possibly other AVs in the intercept network.
- Any other AV can fail at any moment.

Scenarios

Main scenario

1. The AV detects the mechanical fault.
2. The faulty AV communicates its state to the others in the network and to the environment.
3. Faulty AV awaits until it is removed.
4. Before leaving, communicate to the others (and the environment) the resolution.

Postconditions
The network does not contain the faulty AV anymore.

2.1.3 Vehicle software failure

Brief Description

When an Autonomous Vehicle (AV) has a software failure. This implies a mechanical failure, because we assume that a software failure stops the AV. The use case begins with the fault having happened, and ends with the recovery of a normal situation.

Actors

- Autonomous Vehicle (AV)
- Environment

Preconditions

- There is at least one AV with a software fault.
- There are possibly other AVs in the intersection network.
- Any other AV can fail at any moment.

Scenarios

Main scenario

1. One of the other AVs in the network detects the fault.
2. The AV which detects the fault communicates to others in the network which AV has just faulted.
3. Everyone awaits for the faulty AV to be removed.
4. The Environment communicates to the AVs in the network the occured removal.

Alternative scenario

1. On 1 can happen that there is no other AVs, in that case the faulted AV stays into the intersection until someone comes to the rescue.

Postconditions

The network does not contain the faulty AV anymore.

2.2 Assumptions

Some general considerations are here presented. First of all it is assumed a situation in which each autonomous vehicle knows its destination and the roads it is going to travel, since we can assume that each autonomous vehicle has a GPS device on board.

Moreover it is assumed, for sake of simplicity, that all vehicles have the same dimension — or better: that each vehicle can fit into a single position of the internal model used to represent the roads. Moreover, common physical quantities (like weight, speed, acceleration, etc...) are omitted. Instead, the

Figure 2.1: Example of predefined movements that a vehicle can follow in crossing an intersection.

autonomous vehicles move in unit steps governed by the internal model.

Further assumptions are:

- *Faults*: at any moment a failure can arise in vehicles — software or mechanical. A mechanical failure does not compromise the software abilities, but a software failure implies a mechanical failure. So we can assume that if the software fails, the autonomous vehicle stops and goes in "emergency mode". It is also assumed that if a crash happens, it must be managed by removing the faulty vehicle with a tow truck (or something similar).

- *Moves*: the path that a vehicle can take inside the interception is predefined and known a priori. From each position there is a unique path to each destination, and no one can modify or choose a different path once decided. A visual representation of this assumption can be found in Figure 2.1, in which green dots represent positions occupied while crossing, and light-blue dots are the destinations. Once a vehicle reaches its destination, it is considered "satisfied", and can not fail anymore, it can only move forward on its own way. Reversing is not allowed.

- *Message Delivery*: each vehicle can communicate with its immediate neighbors, i.e. the vehicle in front, the one behind it, and if it is at the verge of the crossroads, it can communicate with the others at the verge, and with crossing vehicles. Therefore it's assumed that there exist an upper bound time in message passing between two vehicles. This is justified by the fact that all connection are almost direct, without routers or broker that can cause bottle-necks or network congestion.

2.3 Functional requirements

The solution provides two main modules: *vehicle* and *environment*. Also a *vehicle generator* is provided, in order to test and validate the solution proposed.

2.3.1 Vehicle

Each vehicle has an internal state at any moment. Moreover it:

- knows its position and what happens around, through the environment (in a real context this is handled by sensors like GPS and proximity sensors): it can detect if there are other vehicles trying to cross, or if there is a vehicle in front or behind it.

- communicates with its neighbors by sharing: the direction it wants to travel, its internal state, the next position he is about to reach.

- if it is in the queue, it knows nothing else than its internal information and that it is taking part in a queue with someone before and possibly behind. It can communicate only with these two.

- if it is at the head of the queue, it can connect with other vehicles (that are at the head of other queues) that are crossing the intersection.

So, main functions offered by this module are:

- `startup()` to start the vehicle.

- `stop()` to stop the vehicle.

- `cause_mechanical_failure()` and `cause_software_failure()`, to cause failures; these are useful to test simulating failures.

- `send_message(To,Msg)` to communicate with other vehicles.

- `move()` to move one step forward.

- `check_next_position()` to check if the position it is willing to move is free.

2.3.2 Environment

The environment represent the vehicle ability to use its sensors. It knows the intersection shape and dimension; it knows vehicles approaching the crossroads and the ones in the queues; it provides vehicles with all the information they need to safely circulate within the environment.

So, main functions offered by this module are:

- `is_position_free(Pos)` to simulate proximity sensors; check if a specific position is free.

- `update_position(Vehicle,OldPos,NewPos)` to update its internal state moving the vehicle from the old position into the new one.

- `get_route(StartPos,DestPos)` to simulate the GPS; a vehicle that is going to start its journey ask the environment which is the route it has to travel.

- `get_participants()` to simulate proximity sensors, cameras, and antennas; to ask the identity of vehicles at the verge of the intersection.

2.3.3   Generator

The generator aim is to generate pseudo-randomly an amount of vehicles and start them to test the solution proposed. So it offers a unique function `generate([Prams])`, with some tunable parameters to manage the vehicle generation, such as: how many vehicle it has to generate, and the percentage of failures it has to cause.

2.4   Non functional requirements

The system does not handle byzantine processes nor cybersecurity issues.

- *Safeness*: there can not be more than a vehicle in the same position at the same time;

- *Liveness*: if a vehicle approaches the intersection and is waiting to cross it, eventually it will cross it;

- *Fairness*: if a vehicle approaches the intersection and is waiting to cross it, there exists a bound to the waiting time;

- *Fault Tolerant*: the system is tolerant to mechanical and software failures;

- *Distributed*: there is no central server, vehicles have to coordinate each other.

Chapter 3

Project

This chapter is devoted to the description of the general architecture, and specific algorithms.

3.1 Logical architecture

The solution relies on two main modules, and a secondary one for tests.

- *Vehicle*: it is represented as a finite state automaton. A graphical representation can be seen in Figure 3.1. Each state represents a situation in which the vehicle can find itself. A vehicle is in the *Init* state when it has just approached the intersection. It is in the *Ready* state when it is at the verge and has identified the participants. It is in *Election* when it is performing the election algorithm to choose who will be the one to cross. *Crossing* when it is crossing the interception, and *Terminate* when it has finished the crossing. When a mechanical failure happens, it notifies the event to its neighbors and then terminate; so a crash state is not needed.

- *Environment*: it represents the physical environment accessible through autonomous vehicle sensors.

- *Vehicle Generator*: it is a module aimed to (pseudo-)randomly generate vehicles that approach the intersection, and it is used to test and validate the solution proposed.

3.2 Protocols and algorithms

This section describes the algorithm to solve the core part of the project. Moreover, some peculiar situation sequence diagrams are presented.
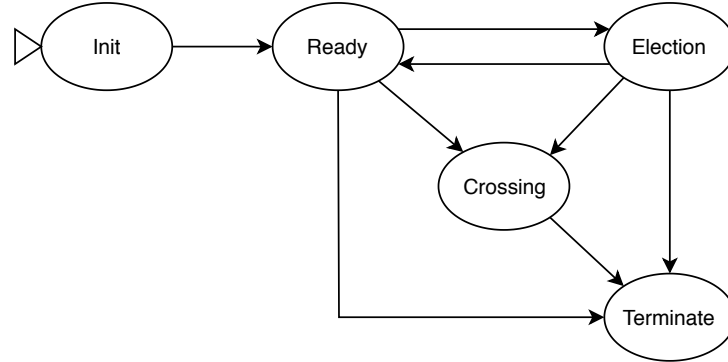
Figure 3.1: Vehicle as a finite state automaton.

3.2.1   Intersection Crossing Algorithm

Vehicles queue up along the roads leading to an intersection awaiting for their turn to cross. The first vehicle of a queue is the vehicle at the verge of entering the intersection and we will call such vehicle a *participant* of the intersection crossing algorithm.

Note that this algorithm applies to generic intersections, just by adapting the representing graph. Moreover, it trivially guarantees the *fairness* and *liveness* requirements (thus it is also *deadlock-free*), thanks to its round robin characteristic.

The algorithm maintains the following invariant through execution:

- one working (i.e. without faults) vehicle per time can be crossing the intersection

- only the leader vehicle can cross the intersection

Algorithm description:

1. The participants (first vehicles in their respective queues) are willing to cross the intersection;

2. They start a slightly modified version [1] of the Bully Algorithm to elect a leader: it terminates with the leader L, and the next leader L', which is the first vehicle after him in a clockwise manner;

3. L begins to cross the intersection;

---

[1] In the canonical version of the Bully Algorithm, every participant is supposed to know other IDs a priori. The version used here does not rely on this assumption. Still, it guarantees the same outcome, even w.r.t. non functional requirements (safety and liveness).

4. After L has successfully crossed the intersection, the participants are informed;

5. L' identifies the next leader L'' by choosing the first vehicle after him in a clockwise manner;

6. The lead passes to L' informing every participant that it's the new leader;

7. The algorithm repeats from 3, where $L = L'$ and $L' = L''$.

Additional details:

- After the leader election, every participant knows who is the current leader, but only the leader knows who is the next one. In this way if the leader dies, it is necessary to perform a new election;

- Once the leader starts crossing, the vehicle behind it (if any) becomes a new participant and asks to join the algorithm. Only the leader answers, providing its identity.

- If a vehicle approaches an empty intersection it automatically elect itself as leader and starts the crossing.

A leader election starts only when new participants do not receive an answer when asking to join the algorithm; or when, after the leader is gone (after completing the crossing, or due to a failure), remaining participants does not receive the message from the candidate, communicating it is the new leader.

3.2.1.1 Managing abnormal cases

All participants monitor only the leader, if it fails, everyone is expecting a message from the candidate, telling it is the new leader. If this not happens, a new election is started. Moreover:

- If the leader fails, the vehicle in clockwise order after him becomes the new leader — identifying also the next candidate leader — and the algorithm restarts from 3.

- If the leader fails while crossing, a timeout T is needed before it is removed from the intersection. Meanwhile a new election can start, and the new leader can start crossing, but it has to stop if the position in front of him is still occupied from the previous faulty leader, until the position has been freed.

- Some peculiar situations can happen for example when a new vehicle arrives at the interception while the leader is gone and it is passing the leading to the candidate. It the new vehicle arrives at an unfortunate moment, it could receive no answers when asking who is the leader, and so starting a new election. The risk is that the candidate can become leader, and after a while someone else can win the election, stealing its

role of leader. This can lead to disastrous scenarios, in which more than a vehicle start crossing. To avoid this, it is *necessary* that a candidate receiving the role from the previous leader, in case of a sudden election, it turns out to be the winner. So, no one can steal the leading from anyone else. This is guaranteed with an appropriate implementation, described in the next chapter.

3.2.2   Sequence Diagrams

This section presents some sequence diagrams of relevant case situations.

3.2.2.1   Normal execution of Intersection Crossing Algorithm

The diagram in Figure 3.2 represent the normal execution of the algorithm described above. It starts at the end of the Bully Algorithm, with the AV_4 communicating it is the leader. It is assumed the queues are enumerated from 1 to 4 in clockwise order.

The leader crosses the intersection by asking the environment if the position in front of it is free, and continues to cross until it reaches its destination.

It is important to note that when the leader starts crossing, AV_5 — which is the one following AV_4 in the queue — moves forward. Then AV_5 is at the verge of the intersection and can communicate with the other participants in order to compete for the crossing.

3.2.3   Dealing with AV failures

When the leader dies because of a software failure, it can not promote the candidate. So a new election is needed. The situation is represented by the sequence diagram in Figure 3.3

The last situation in Figure 3.4 represents a failure of the candidate leader. When the leader finish the crossing, everyone wait a timeout T for a message from the candidate. Since it is not alive, the timeout runs out, and a new election is performed.

3.3   Physical architecture and deployment

The architecture is quite simple, a graphical representation is displayed in Figure 3.5. Each Autonomous Vehicle (AV) is a physical node. The Environment represents the world sensed by AVs, so it is on a separate node, together with the vehicle generator. AVs can ask the environment for information, and they can communicate with their neighbors [2].

---

[2]For a definition of "neighbors", see in Chapter 2, in assumptions section. Where the message delivery assumption is described .
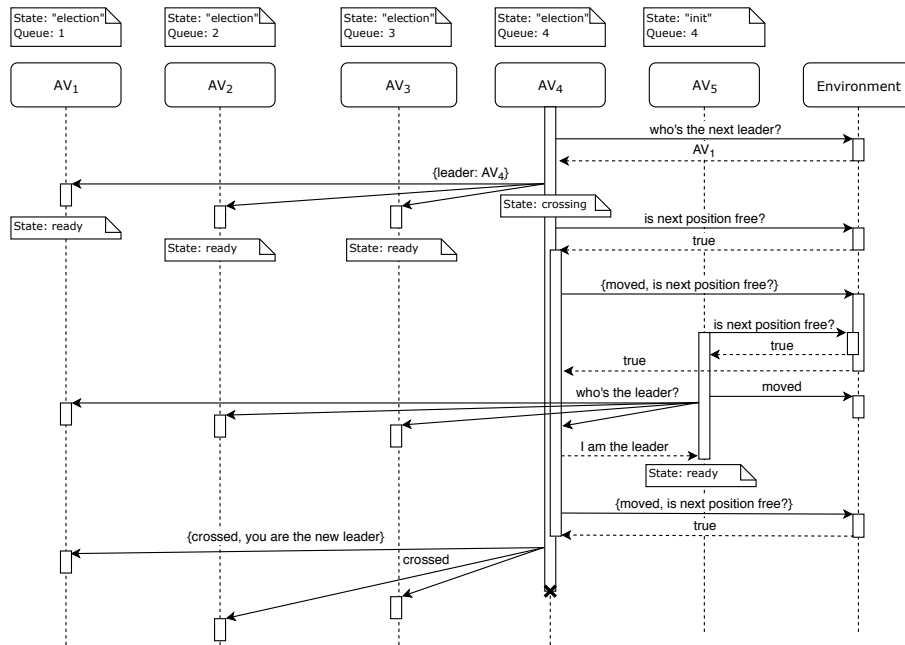
Figure 3.2: Sequence diagram on a normal execution of the Intersection Crossing Algorithm described above.
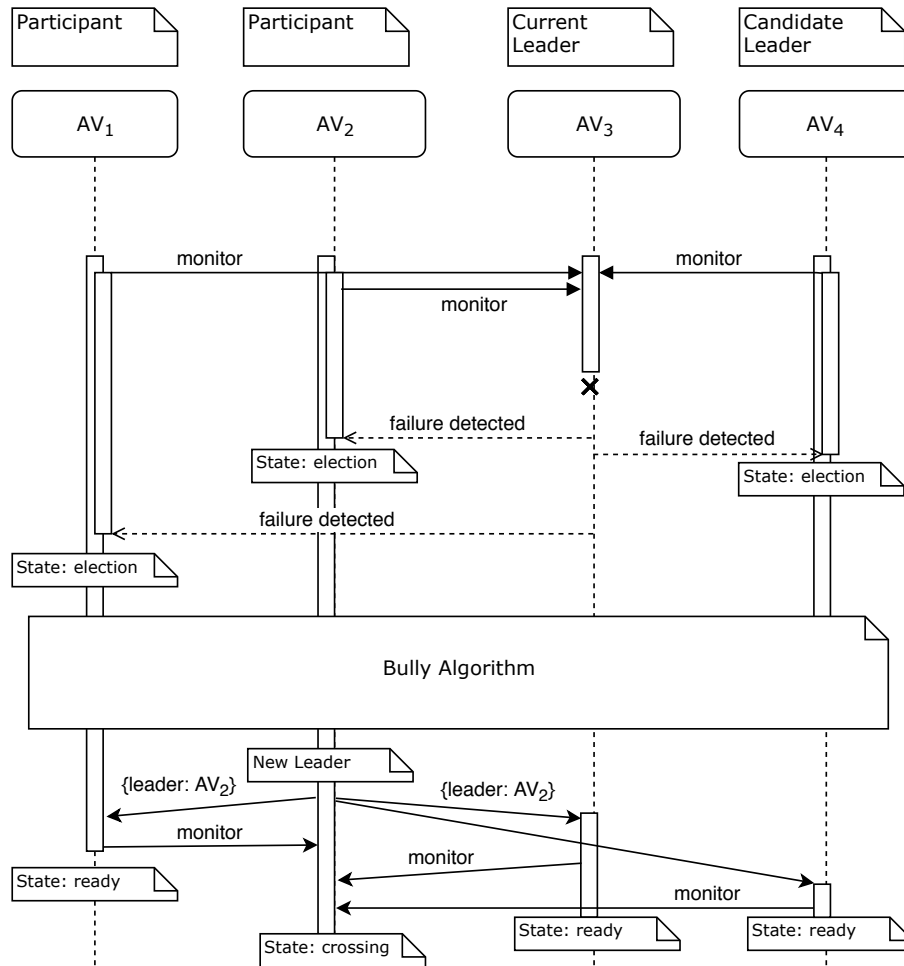
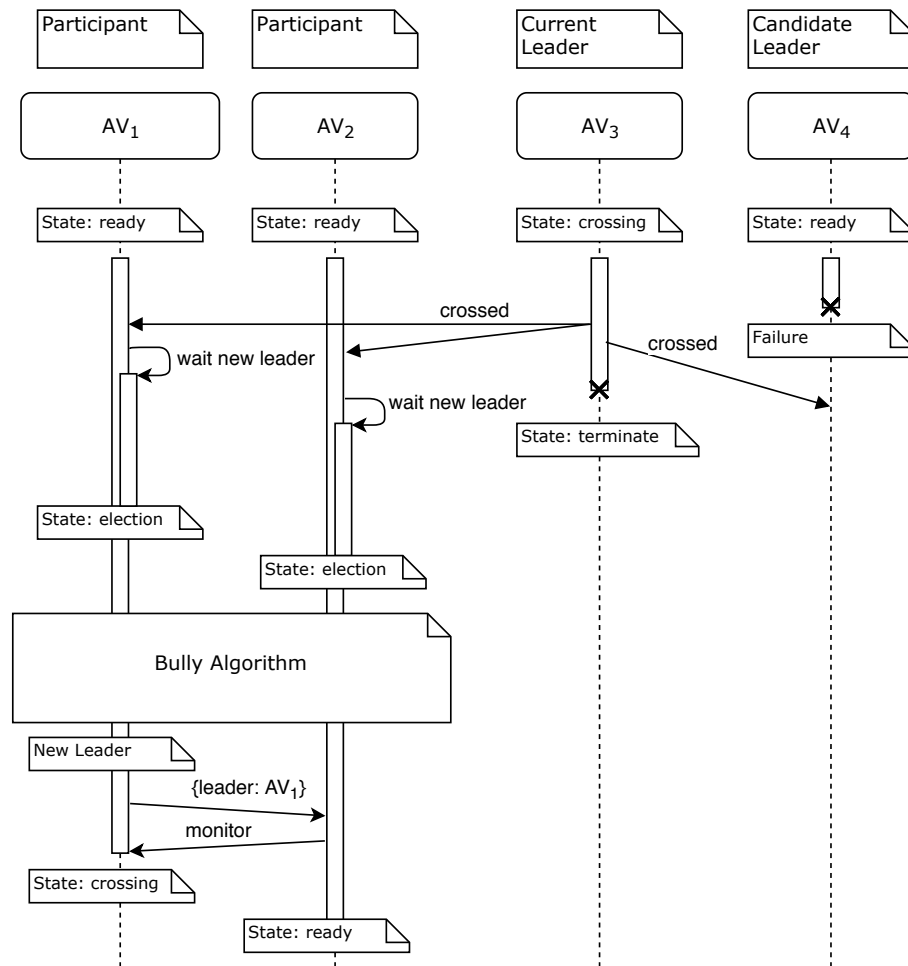Figure 3.3: Sequence diagram representing a generic leader failure.

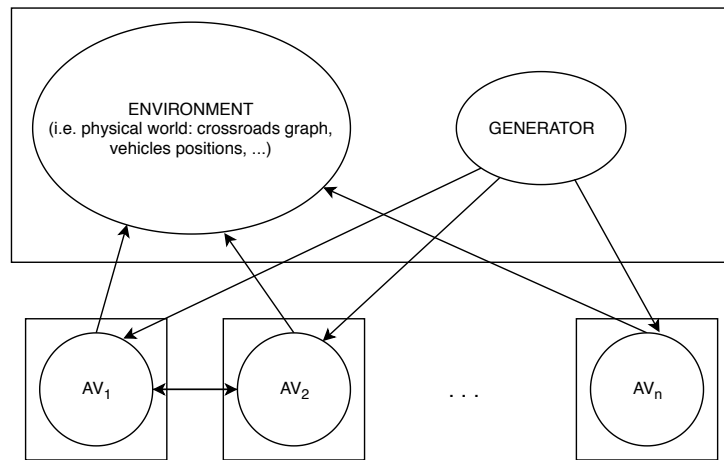Figure 3.4: Sequence diagram of a candidate leader failure.

Figure 3.5: Physical architecture: Environment, Generator, and Autonomous Vehicles (AV).

Chapter 4

Implementation

The project is developed in Erlang. For the development we used as much as possible OPTs offered by Erlang, since they speed up the work and offer a solid and well-tested starting base.

4.1 Getting Started

4.1.1 Project Structure

The project is organized in two different Erlang application: Vehicle and Environment. These can be found in the `apps/` directory. Then there is the `test/` directory which contains some scripts to start the tests (described in Chapter 5), and the vehicle generator module. Last but not least there is the `docker/` folder and the `Dockerfile`, that are used to test the project with Docker (even this is described in Chapter 5).

4.1.2 Hands On

To handle the whole project is provided a *makefile*, which lets to compile and clean the project, and provide also a couple of demos.

Make commands are described here:

- `make demo` performs a demo with erlang nodes generated in local host.

- `make dockerdemo` performs a demo using Docker. Note that this commad will create some containers and a docker network on your computer. If you want to clean up everything later see `make clean`.

- `make` (which is the same as `make apps`) compiles just the apps: environment and vehicl.

- `make all` does what `make` do, but it compiles also the generator module and gives the scripts the execution permissions.

Figure 4.1: Graphical representation of the intersection graph contained in the configuration file, and used to test the project. Nodes with name R_x represents roads. Details: in green start nodes (R_1, R_11, R_21, R_31), in blue destination node (R_10, R_20, R_30, R_40), in yellow intersection internal nodes (I_x), in red intersection entrance nodes (IN_x), in purple intersection exit nodes (OUT_x).
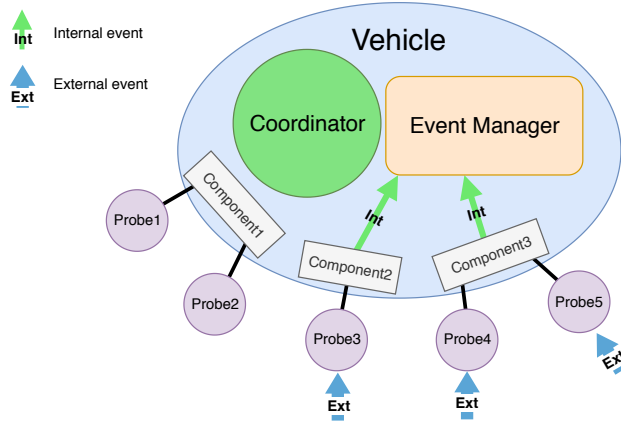
Figure 4.2: Vehicle system architecture: Coordinator, Event Manager, and Components with Probes attached.

- `make docker` does what `make all` do, but it compiles also the Docker stuff.

- `make clean` removes all compiled files, log directories, and Docker stuff.

- `make distclean` removes the same as `make clean`, plus it cleans also the report directory which contains latex files.

4.2 Development Details: The Vehicle

4.2.1 Vehicle system

The vehicle system architecture follows the event-based multi-agent architecture where the only source of events are vehicles (agents) and the environment. Such architecture will greatly help in dealing with accidental complexity, changing requirements, security and performance.

Vehicles represent the main source of events, autonomously driving control towards their own goals, and producing internal events through their actions. The environment and other vehicles model the external events. External events are captured by the vehicle's devices such as sensors, sonars, camera systems, receivers, transmitters, etc.

The main elements of the architecture are the Coordinator, the Event Manager and various components that provide a specific service. Vehicle components communicate by producing and receiving event notifications. All of the vehicle's functionality is split among the components: each component is responsible for delivering a specific service to the vehicle.

Components can be connected to sensors and other external devices in order to successfully deliver their functionality. External devices are referred to as
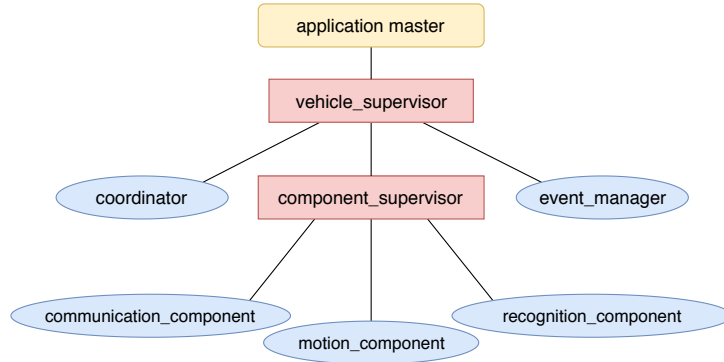
Figure 4.3: The implemented vehicle system supervision tree (Erlang point of view).

probes. The responsibility of probes is to collect external events from the environment or other vehicles and forward them to the component to which they are attached to. External events received by components and translated them into internal events.

The Coordinator listens to specific internal events in order to coordinate the components to reach the vehicle's goal.

4.2.2   Architecture elements

This subsection presents the architectural elements: Event Manager, Components and the Coordinator. Figure 4.3 shows the supervision tree implemented in Erlang with all the architectural elements.

4.2.2.1   Event Manager

Components communicate by sending event notifications to the Event Manager. Being not real messages, events have not a designed receiver: they are distributed by the Event Manager to Components that subscribed to that sort of events.

The Event Manager acts as a middleware, inverting the logic of program execution, where a component does not need to know the existence of other components to provide its service and thus facilitatates the addition of new features and components.

4.2.2.2   Components

Components are the only providers of functionality. Each component delivers a specific functionality to the system.

Components can be *passive* or *active*. Passive components do not generate internal events and their sole purpose is to listen to internal events and act

accordingly. An example of a passive component is the Logging Component which listens for specific internal events and logs them onto a file.

Active components are components that generate internal events. For example the Communication Component, whose sole purpose is to handle communication between vehicles, is required by components that need Vehicle-to-Vehicle (V2V) communication services.

Note that both active and passive components are capable of performing actions, changing their state and communicating with their devices except that passive components cannot interfere with the event flow of the vehicle because they don't generate internal events.

Components can have multiple devices attached (sensors, receivers, transmitters, etc.) in order to detect external events and provide their functionality.

Components help to address the issues of event aggregation and transformation by making event notifications meaningful at the level of interpretation required for making other components activities effective. They effectively act as event mediators (or, correlators), translating external events into meaningful internal events.

Component Structure

The structure of a component consists of:

- component module (mandatory)

- event listener (optional)

The component module:
The component module maintains all the logic related to the specific service it is built for, for example a component in charge of the navigation system deals with navigation related tasks only.

The component module can be connected to zero or more devices in order to provide its services, for example the navigation component can use a GPS device to gather geolocation data.

When needed, components can spawn additional processes called Actions to perform more advanced and complicated tasks (Actions can also be used for enabling the selection of algorithms at runtime, effectively letting the algorithm vary independently from the component, a strategy pattern in OO terms).

Picture 4.4 shows a component with its actions in the supervision tree. Actions are attached to a specific supervisor (the Action Supervisor) which is unique for each component.
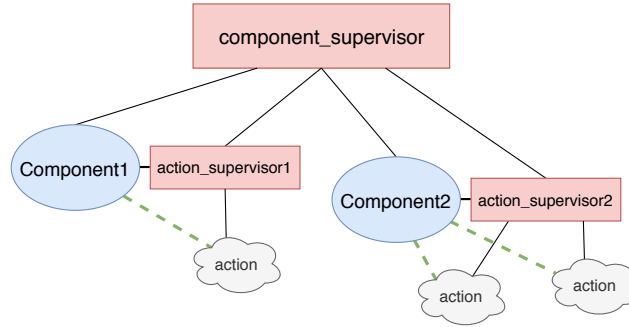
The event Listener:

Figure 4.4: Part of the overall supervision tree showing components with their actions and their relative Action Supervisor


Figure 4.5: Components with Event Listeners registered on the Event Manager

A component module may have zero or more event listeners attached to the Event Manager. Event listeners are responsible for listening and forwarding internal events that are of interest to the component module.

A component that provides a service must have an event listener in order to receive requests from other components. Event listeners are event handlers in Erlang terminology.

Figure 4.5 shows how event listeners are attached to components and registered on the Event Manager. Note that Component4 doesn't have any event listener, which is a perfectly valid situation.

4.2.2.3   Coordinator

The Coordinator represents the logic and reasoning of the system, coordinating the various components in order to achieve the vehicle's goals.

The Coordinator listens for internal events of interest by registering event listeners on the Event Manager (just like components do). Specific internal events trigger the Coordinator's actions which, in turn, produce new internal events addressed to specific components.

This behaviour of recognising specific events and addressing specific components is what drives the vehicle towards its goals.

4.2.2.4 The general flow of events

Once the vehicle has been initialized and the startup message received the vehicle can begin to move.

(1)The Coordinator starts by asking the Recognition Component about the position type it is currently on.

There can be 2 basic position types (although the implementation provides significantly more) a *free* position and a *stop* position.

A *free* position is any position onto which the vehicle is free to advance (for example, it is on a public road), where as a *stop* position is a position that requires the vehicle to stop and resolve the situation (for example, an intersection entrance, an obstacle on the road, a car accident in front, etc).

The Recognition Component utilises, for example, a camera system to detect the position type and returns the answer to the Coordinator.

If the answer is *free*, the Coordinator requests the Motion Component to check the position in front whether the vehicle is safe to advance. If the answer is *stop* the vehicle has to stop and resolve the problem. For the sake of simplicity (and, for the scope of this project) we assume that a position has type *stop* if the vehicle is at the entrance to an intersection.

(2)Suppose that the Recognition Component's answer was *free*, the Coordinator requests the Motion Component to verify the position in front. The Motion Component uses a proximity sensor to check whether the position in front is clear or occupied. The answer is provided to the Coordinator.

If the answer is *clear*, the vehicle advances in the next position. If the answer is *occupied* the vehicle has to realise what is in front. For the scope of this project, we assume that only vehicles can be on the road. Therefore the vehicle tries to communicate with the vehicle in front, in order to check its health. The Coordinator requests the Communication Component to contact the vehicle in front.

The Communication Component can use, for example, a Bluetooth device or a Wifi device to communicate with the vehicle in front (for example, send a broadcast with a message: "Who is at position X?"). The vehicle in front can

either respond to the call, or, if it has software crashed, not respond at all. An answer is provided to the Coordinator.

The Coordinator can either receive that the vehicle in front is *up* (software is running correctly and the vehicle is waiting in a queue) or *down* (the vehicle in front has software crashed). If the vehicle in front is *up*, the vehicle awaits in queue for the vehicle in front to move. If the vehicle in front is *down* the vehicle calls a tow truck to remove the vehicle in front and awaits its removal. In either cases, after a period of waiting, the vehicle moves in the position in front and the algorithm starts from (1).

Returning to the step when the Coordinator requests the position type. If the Recognition Component's answer is *stop*, the vehicle is at the entrance to an intersection (for the reasons stated above) and has to stop to resolve the intersection crossing with vehicles at the other entrances in order to avoid potential collisions.

The Coordinator uses the Communication Component in order to solve the intersection. The Communication Component contacts the intersection participants (vehicles at the entrances) and starts the intersection coordination (described in the next section). Once the vehicle is clear to move in the intersection, a notification is made to the Coordinator that the current position is *free* and the algorithm takes on from (2).

4.2.3   Intersection Coordination

This module is the core of the crossing. Here is implemented the algorithm described in Section 3.2.1. This module is an OTP `gen_statem`.

The main idea is to use as ID the couple (`waiting_counter, nonce`), where the nonce is a unique code for each vehicle (the Erlang node name), and the waiting counter is a counter that is zero when a vehicle approach the interception, and increases by one for each `crossed` message it receives, i.e. every time the leader cross, the participants increase their counter. In this way, every time an election is performed, the priority goes to the ones that waited longer.

Moreover, is important to guarantee that no one steal the leader title to a candidate who received promotion from the previous leader — as specified in Subsection 3.2.1.1. To do so, when a leader pass the lead to the candidate, the candidate set its wait counter to N, where we define N as the number of interception entrances. This guarantees what stated before, because no one can wait more than N, or better: every vehicle waiting counter reach at most $N-1$. So, if an election is performed, the candidate will win for sure. Setting the counter to less than this can lead to problems.

Last important part is that since no one knows others IDs, it is necessary to exchange that information while performing the election, so the algorithm is developed as follows:

- to start the election send election message with my ID, and start waiting for answers;

- if an answer is received then this means there exists someone with an higher ID, and I start to wait for the coordinator message;

- if an election message is received, see the ID inside, and compare it with mine. If the ID received is higher I start waiting for a coordinator message, otherwise I forward the election message to all participants, with my ID inside, and then, start waiting for answers;

- if I don't receive any answer message, I understand I have the maximum ID, so I am the leader and I sen the coordinator message.

Some timeouts are necessary to develop this idea. All the details can be found in the code.

4.3   Development Details: Environment and Vehicle Generator

The Environment is an application where its core part is an OTP `gen_server`, which act as a substitute of the vehicle sensors. It provides mainly the functions described in the functional requirements analysis.

The interception graph is provided in a configuration file inside the `config/` directory. The syntax is trivial, so starting from the configuration file provided is easy to build a different graph. Note that are not performed checks on the graph structure, so it is up to the user to provide a correct configuration file.

To manage the graph, the *digraph* Erlang module is used. In particular the route from the starting position until the destination is the shortest path between the two nodes in the graph.

The generator is a simple OTP `gen_server`, which offers only the function `start()`, with different tunable parameters. Details of this module are described in next chapter.

Chapter 5

Validation

The validation part consisted in three main phases. First of all the project was tested with few vehicles generated singularly into different shells, with Erlang nodes created on local machine. After the initial validation, simulating peculiar situation manually (software and mechanical failures, vehicles arrival at specific times, . . .), the software has undergone numerous tests with the help of the vehicle generator described in the previous chapter, even this with Erlang nodes on local machine. Finally, exploiting Docker facilities, we tested the project even in distributed containers on different hosts in a virtual network.

The project has been tested on Fedora 30 with Erlang/OTP 21, and on Ubuntu 18.04 LTS with Erlang/OTP 22.

After careful planning and developing, and after all these tests phases, we are confident of providing a reliable and sound solution. In next sections is described how to reproduce the tests. All the scripts provided have been tested only on cited above operative systems, same results are not guaranteed in Mac OS or Windows.

To perform tests is necessary to have in mind the graph of the interception presented in Figure 4.1. Moreover, for the test with the generator and with docker, the scripts will not exit at the end of the simulation, because we can't understand when all vehicles have finished, but you can easily check that if for about a minute nothing is printed, then the simulation is over, and you can kill the script with `ctrl+C` and then `(a)bort`.

5.1   Simple Tests

To use the code proposed here is necessary to execute `make all`. Some scripts are provided in the direcotry `test/`. Note that is is important to execute the script provided from the project root.

First of all it is necessary to start the environment. To do it open a shell and run:

```
./test/start_environment.sh
```

You should se now an Erlang shell with the message "Env started!".

Now it is time to produce vehicles: pretty easy, open another shell and run:

```
./test/start_vehicle.sh 1 R_1 R_30
```

Where, the first parameter is the index of the vehicle to generate — note that if you insert x the erlang node will be vx@hostname. The second and the third parameters are the start position, and the destination position, chosen from the graph in Figure 4.1. One can also start a vehicle from other positions chosen from R_x or IN_x, to test different situations.

You will see that in every vehicle shell is printed what it is doing, and more information can be found in the vehicle logs inside the log/ directory.

5.2   Vehicle Generator

To use the code proposed here is necessary to execute make all. To start the vehicle generator can be used the script inside the test/ directory, as follows:

```
./test/start_generator.sh <vehicle number> <fail ratio> \
[relaive sw fail ratio] [max fail timeout (ms)]
```

for example you can try:

```
./test/start_generator.sh 10 0.2 0.5 20000
```

Where:

- <vehicle number> is mandatory and represent the number of vehicles that will be generated.

- <fail ratio> is mandatory and represent the percentage (expressed with a number $0 \leq x \leq 1$) of failures that will be generated.

- [relaive sw fail ratio] is optional and represent the ratio of software failures caused by the generator. (e.g. if set to 0.5, half faults will be software, and the other half will be mechanical).

- [max fail timeout (ms)] represent the maximum time in milliseconds within which the fault will be caused, considering the moment when the vehicle starts. (e.g. if set to 20000, the vehicle can fail after a maximum of $20s$ from its start).

To stop the execution press ctrl+C and then (a)bort. The information here are printed all in the same shell, so it is less clear what is happening, but the logs for each vehicle can be found in the log/ directory.

5.3   Docker

To use the code proposed here is necessary to execute `make docker`. Note that with the tests proposed here is required to have docker installed and configured properly. Some containers, and a network, will be created. In `docker/` directory are provided some script to easily run the tests and also to stop and clean the containers and network created.

The command to run is:

```
./docker/start_docker.sh <vehicle number> <fail ratio> \
[relaive sw fail ratio] [max fail timeout (ms)]
```

for example you can try:

```
./docker/start_docker.sh 10 0.2 0.5 20000
```

Where the parameters are the same of the vehicle generator described in the previous section. Note that in docker you will se only messages printed from the generator. To see what happens to the vehicles you can look at the `docker_log/` directory, where all vehicles log are generated. Further, can be necessary to provide root permission to delete `docker_log/` directory, because the files are generated by root user inside the containers, so the proprietary user is still root. After stopping the simulation, everything can be cleaned with:

```
./docker/clean_docker.sh
```

Chapter 6

Conclusions

The proposed solution focuses on a very simplified version of what a real solution could be. In particular the aim was to produce a core solution which satisfies the requirements but that also does not contain more than necessary assumptions, and in particular that does not rely on non realistic ideas that can lead future development to a blind spot.

6.1 Partial implementation and future work

Some advancement that can be pursued starting from our solution are provided here. Some of them are easier than other, but due to lack of time we did not managed to deepen them:

- Try different interception graphs;

- Centralize some configuration parameters in such a way the solution can be tuned without recompiling;

- Exploit code change feature offered by Erlang;

- Develop a graphical interface of the environment to better understand how vehicles are moving;

- Add vehicles with priorities, e.g. emergency vehicles;

- Dynamically modify paths inside the crossroad if some position becomes unavailable;

- Let more than one vehicle per time cross the intersection, maybe with an interface to a planner in order to find the best route (exploiting epistemic planning could be an interesting option).